

Process, IPC, and Threads

Chapters 3 & 4

Today's Topics

- Complete OS Design Discussion
- C Tutorial/Refresher
- Introduction to Processes
- Creating Processes - POSIX

Dual Mode Execution

- Implemented using Trap/Interrupt and Mode-Bit
 - Must be supported by hardware
- Privileged Instructions cannot be executed by user code
 - Allows OS to make sure programs “play nice”.

Designing an OS

- OS is still quite complex and difficult to write!
- Design impacts both functionality and quality
- Simple Design: Monolithic
 - MS-DOS
 - Original versions of UNIX
 - Why?

Varying Priorities

- Many implementation choices for:
 - Scheduler
 - Memory Management
 - Filesystems
 - Security/Communication
 - etc.
-

Modern Design Principles

- ▣ Layered
 - ▣ Advantages: Layers interchangeable, easy debugging
 - ▣ Disadvantage: Interdependencies
- ▣ Micro-Kernels:
 - ▣ Advantages: Flexible, Secure
 - ▣ Disadvantage: Takes Discipline, Inefficient
- ▣ Modules:
 - ▣ Advantages: Flexible, Secure, Efficient

C Tutorial

- Different includes
 - Ex. `<stdio.h>`, `<stdlib.h>`, `<math.h>`
- Variable Declarations - you must declare variables at the beginning!
- C-Style I/O - File Descriptors
 - `printf`, `fprintf`
 - `fgets`
 - `fgets (char * str, int num, FILE * strm);`
 - `fopen`, `fclose`

C Data Structures

- Declaring a data structure:

- `struct myStructure x;`

- `struct myStructure *x;`

- `x = malloc(sizeof(myStructure));`

- Defining structures:

- ```
typedef struct myStructure {
 int field1;
 int field2;
} myStructure;
```



# Other C “Quirks”

- No native pass by reference
  - use pointers
- no `const`, use `#define`
- POSIX uses structures extensively.

# C-Strings

- POSIX was written long before the C++ STL libraries and other extensions.
  - It will not play nicely with the string datatype.
  - Always use c-strings

|   |   |   |   |   |   |   |   |    |  |  |
|---|---|---|---|---|---|---|---|----|--|--|
| c | - | s | t | r | i | n | g | \0 |  |  |
|---|---|---|---|---|---|---|---|----|--|--|

<http://www.cplusplus.com/reference/clibrary/cstring/>

# Some Rules

- Your programs can be C++, but should **not contain** the following:
  - C++/STL string classes (use cstrings)
  - No STL (vector, map, etc.)
  
- You **can** use
  - C++ or C input/output (i.e. cout)
  - classes (unless otherwise specified)

# Program 1

- C-Demo.c
  - *Examine C-style I/O and Structures*

# Processes

- Process v.s Program
  - A Process is a *running* program, with data, stack, heap, etc.
    - What does this look like in memory?
  
- Process States:
  - new
  - ready
  - running
  - waiting
  - terminated

# Implementation

- Each Process represented by a **Process Control Block**
- Scheduling Queues implemented using linked lists
  - A schedule manages several queues
- Context Switch
  - Requires PCB to be saved/restored
  - What's in the PCB?

# Process Creation/Termination

- Processes are organized in parent/child relationships
- Each process has a unique ID (integer)
- Parent may:
  - Kill off children (implicitly or explicitly)
  - Wait for children to terminate
  - Orphan its children

# POSIX Process Management

- Program can create child process using `fork()` system call.
- Address space of parent is *completely* copied into child's.
- Child begins execution immediately after fork
- Parent resumes execution after fork call
- Different return values



# Understanding Fork

```
int pid;
.
.
pid = fork();
if (pid == 0) {
 // this is the child process
} else {
 // this is the parent process
}
```

One process  
(parent) executing...

After fork, two processes  
executing same code...

# Program 2

## forks.c

*Examine behavior of multiple process programs.*

---

# Endless Forks...

- Symptom:

- "I can't log into my account anymore"
- "phobos has slowed to a crawl"
- "Strange things are happening..."

- Common Causes:

- You are creating too many processes (and they are not terminating)
- You are calling fork in a loop!

# Controlling Child Processes

- **Please be pro-active** - if you end up with too many rogue processes everyone suffers, you get locked out, and the sysadmin needs to fix it.
- When developing your code, always print out process creation/termination information and verify it make sense
- Use the `ps` command to monitor your active processes
- Use the `kill` command to kill off erroneous processes

# Controlling Child Processes

```
~/classes/cms311> ps -fu sfrees
```

| UID    | PID   | PPID  | C | STIME    | TTY    | TIME | CMD            |
|--------|-------|-------|---|----------|--------|------|----------------|
| sfrees | 23874 | 1     | 0 | 10:05:47 | pts/30 | 0:00 | killdemo       |
| sfrees | 23875 | 23287 | 0 | 10:05:56 | pts/30 | 0:00 | ps -fu sfrees  |
| sfrees | 23285 | 23282 | 0 | 09:43:50 | ?      | 0:00 | /usr/lib/ssh/s |
| sfrees | 23287 | 23285 | 0 | 09:43:50 | pts/30 | 0:00 | -bash          |
| sfrees | 22810 | 675   | 0 | 09:28:05 | ?      | 0:00 | /usr/local/sam |
| sfrees | 23534 | 23287 | 0 | 09:51:19 | pts/30 | 0:01 | /opt/csw/bin/e |
| sfrees | 23599 | 23287 | 0 | 09:52:35 | pts/30 | 0:01 | /opt/csw/bin/e |

```
~/classes/cms311> kill 23874
```

```
~/classes/cms311> ps -fu sfrees
```

| UID    | PID   | PPID  | C | STIME    | TTY    | TIME | CMD            |
|--------|-------|-------|---|----------|--------|------|----------------|
| sfrees | 23891 | 23287 | 0 | 10:06:24 | pts/30 | 0:00 | ps -fu sfrees  |
| sfrees | 23285 | 23282 | 0 | 09:43:50 | ?      | 0:00 | /usr/lib/ssh/s |
| sfrees | 23287 | 23285 | 0 | 09:43:50 | pts/30 | 0:00 | -bash          |
| sfrees | 22810 | 675   | 0 | 09:28:05 | ?      | 0:00 | /usr/local/sam |
| sfrees | 23534 | 23287 | 0 | 09:51:19 | pts/30 | 0:01 | /opt/csw/bin/e |
| sfrees | 23599 | 23287 | 0 | 09:52:35 | pts/30 | 0:01 | /opt/csw/bin/e |

```
~/classes/cms311>
```

# Inter-Process Communication

- Two common types:
  - Shared Memory
  - Message Passing
  
- POSIX Pipes
  - Typically implemented using shared memory
  - FIFO, one way communication
  - Also partially supported on Windows

# PIPES

```
int pfd[2]; // file descriptors
```

```
int result = pipe(pfd);
```

*After pipe call pfd[0] contains descriptor for read end of pipe, pfd[1] contains the descriptor for write end.*

Use read and write system calls to read and write pipe. These calls require an fd, which is the file descriptor (pfd[0] for read, pfd[1] for write).

```
int read(int fd, char * buffer, int maxlen)
```

```
int write (int fd, char * buffer, int numBytes);
```

# Program 3

- Pipe-example.c

- *Sending information from one process to another*



# Exec System Call

- Having two processes active in the same code causes confusion:
  - Rarely used.
- Normally when processes are created, they are created to run different programs.
- POSIX: `exec*`
- Win32: `CreateProcess` (Windows does not provide a separate `fork`)

# Program 4

- Echo

- *Inter-process communication with separate programs*

## Next Class:

- Programming with multiple threads
  - POSIX
  - Win32
  
- Contact me if you have not been able to get into [cs.ramapo.edu](http://cs.ramapo.edu) (or a linux or Mac machine)
  
- Homework #2 is assigned
  - Part 1 is processes (start now!)
  - Part 2 and 3 is threads (start after next class)