

Threads

Chapter 4

Processes

- When creating a process, the child process gets its own resources:
 - Address Space (global variables, files, etc)
 - heap
 - “Thread” of execution
 - Program Counter
 - Stack
- This can be a good or bad thing...

Inter-Process Communication

- Two common types:
 - Shared Memory
 - Message Passing

- POSIX Pipes
 - Typically implemented using shared memory
 - FIFO, one way communication
 - Also partially supported on Windows

PIPES

```
int pfd[2]; // file descriptors
```

```
int result = pipe(pfd);
```

After pipe call pfd[0] contains descriptor for read end of pipe, pfd[1] contains the descriptor for write end.

Use read and write system calls to read and write pipe. These calls require an fd, which is the file descriptor (pfd[0] for read, pfd[1] for write).

```
int read(int fd, char * buffer, int maxlen)
```

```
int write (int fd, char * buffer, int numBytes);
```

Exec System Call

- Having two processes active in the same code causes confusion:
 - Rarely used.
- Normally when processes are created, they are created to run different programs.
- POSIX: `exec*`
- Win32: `CreateProcess` (Windows does not provide a separate `fork`)

Program 4

- Echo

- *Inter-process communication with separate programs*

Threads

- Can be considered *lightweight* alternative to using multiple processes
- Threads **share** global address space, heap, and code
- User v.s Kernel Threads
 - Pthreads, Win32, Java

Thread Resources

- One process can contain many threads
 - Each thread has its own execution context
 - All threads share the same address space and resources
- Advantages: Easier to work together
- Disadvantages: OS does not provide protection among common threads!

Scheduling Threads

- We know an OS scheduler switches between “ready” processes. How does it determine which thread to run?

Depends on the implementation:

Kernel Threads v.s. User Threads

User Threads

- User Threads: The OS knows nothing
 - Implemented by a runtime system, linked to the program code itself.
 - This runtime is almost like a tiny OS!
 - Used to be common (dominant), no longer.
 - Scheduling is not pre-emptive.
 - Major advantage: Extremely fast.
 - Major disadvantages:
 - Multi-processor?

Kernel Threads

- Kernel threads are managed by OS:
 - OS contains thread tables, indicating how many runnable threads are in the process
 - OS Schedules not only process, but choses among threads
 - Switching threads is more costly (why?)
 - Advantage:
 - Multiple threads can run on separate processors
 - One thread blocks on read, others can continue processing.

Why use threads?

- We need to remember programs usually cycle between CPU use and I/O use
 - The overall system is better utilized if CPU **AND** I/O are always busy
 - Your program is optimized if it can use the CPU **and** I/O simultaneously!
- This optimization is only available when using kernel threads (**why?**)

Classic Programming Models

- Multi-Threaded Web Server
- Multi-Threaded User Interface
- Multi-Process distributed system
 - Multiple machines may/may not be involved

More considerations

- Threads are faster to start/create than processes
 - Very fast for User Threads
 - Fast for Kernel Threads (less is copied)
- Thread Pools
- Termination
- Synchronization and Data Protection
 - Chapter 6

First: Review of Function Ptrs

- In both POSIX and Win32, creating a thread requires you to provide a function pointer
 - CreateThread calls this function (with whatever parameters you give it) as soon as it creates the new thread
 - Just like variables, functions have “types”
 - Type equivalent to function *signature*
 - The function’s name is a pointer to “code” instead of data

Function Pointers

```
// Two function, both with the same signature
```

```
int add(int a, int b);
```

```
int sub(int a, int b);
```

```
void print(int a, int b);
```

```
// Function that takes a and b and performs the operation
```

```
// by executing the function provided:
```

```
int execute(int a, int b, int (*function)(int, int) ) {
```

```
    return function(a, b);
```

```
}
```

```
// a call to the execute function:
```

```
execute (7, 6, sub);
```

```
execute(7, 6, print); // print doesn't have the required signature
```


Thread Libraries

- pthreads: POSIX standard
 - C procedure interface
 - Kernel or User depending on underlying OS
- Win32 Thread Model:
 - Kernel threads

Pthreads

- Must include `pthread.h`
- Pthreads start in a user defined function with specific signature:
 - `void * function_name(void * param)`
- Functions:
 - `pthread_attr_init` (set thread attributes to defaults)
 - `pthread_create` (create thread)
 - `pthread_join(id)` (wait for thread (id) to terminate)
- **compile with `-lpthread` option**
 - `g++ -o myprog myprog.cpp -lpthread`

Program

■ Pthreads

Win32 Thread API

- Windows uses HANDLES to represent system-wide identifiers:
 - processes, threads, files, etc.
- HANDLE CreateThread(...)
 - security attributes (NULL)
 - default Stack Size (0)
 - thread function
 - parameters to thread function
 - creation flags (0)
 - &thread identifier (DWORD, not really used)

Win32 API

- Thread function has strange signature:
 - `DWORD WINAPI functionName(LPVOID parameters)`
- To wait for a thread to terminate, you specify the handle:
 - `WaitForSingleObject(ThreadHandled, INFINITE)`
 - You can also specify milliseconds to wait...
 - After thread has terminate, use `CloseHandle()` to clear up resources

Program

- Simple Win32 example
- Thread Summation Example

Threads in Java

- Java also supports threads - in fact, quite elegantly
 - Objects can implement the *Runnable* interface
 - required to have a void run() function
 - The object can be run as a separate thread (starting in the run function)
 - Member variables are shared.

Complications

- Do Processes inherit all threads?
 - Only an issue with POSIX, since Win32 always executes a new program image
 - What thread receives keyboard / mouse input?
 - What thread receives any OS signal?
- Typical Solution: Its up to the programmer...

Summary

■ Multi-Process Programs:

- Increases parallelism
- Good when a job can be broken into multiple independent tasks

■ Multi-Threaded Programs:

- Increases parallelism
- Good when job can be broken into multiple complimentary tasks
- Multi-Threaded programs are extremely common and simplify design of complex systems

Exam 1

Exam Details

- Exam is closed book / closed notes
- Covers Chapters 1 - 4
- You will have the full class period for the exam
 - Short Answer/Multiple Choice
 - Longer Problems
 - You are **not** required to memorize function names/parameters
 - Any code you are asked to write will be psuedo-code

Responsibilities of an OS

- CPU Allocation
- Memory and Storage management
- Regulate and Provide Access to
Peripherals

Defining/Characterizing an OS

- ▣ Services provided to the user ***
- ▣ Architecture
- ▣ System Calls

System Calls

- OS exposes functionality to applications via **system calls**
 - From a programmer's perspective, they are nothing special...
 - Their implementation is **very** different however...
- *Understand dual-mode execution and interrupt handling*

Processes

- Process v.s Program
 - A Process is a *running* program, with data, stack, heap, etc.
- Process can be in 5 different states
- OS represents each process by a **P**rocess **C**ontrol **B**lock
- *Understand what fork and exec calls do...*

IPC

- Inter-process Communication:
 - Shared memory
 - Message Passing
- Understand pipes and how they can be implemented

Threads

- Threads consist of separate stack, registers, program counter
 - All threads within a process share code and heap
- Kernel v.s. User Threads
- Understand thread creation API calls