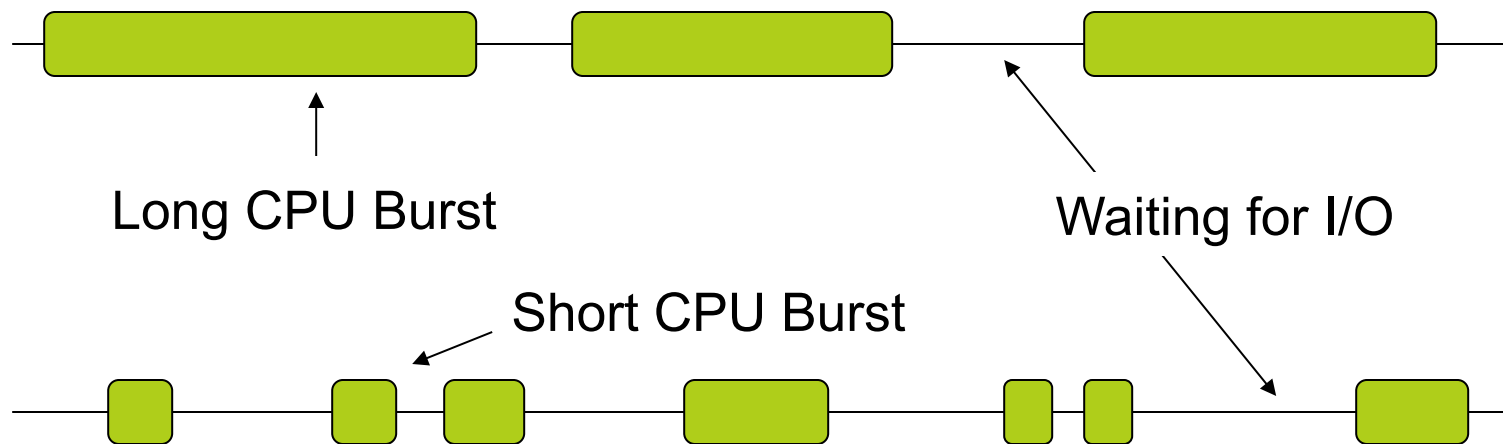

Scheduling Algorithms

Chapters 5

CPU Bursts



Ideally, we want the CPU **and** disk to be *always* busy.

CPU-Bound: Characterized by long CPU bursts

I/O-Bound: Characterized by short CPU bursts

Scheduling Decision Points

- Standard Approaches:

- Process Creation
- Process Termination
- Blocking Call is issued
- Process Yields voluntarily

- Pre-emptive

- Blocking Call completes
 - Timer interrupt
-

Scheduling Algorithms

- Scheduling Criteria (Goals)
 - different goals for different systems
 - Batch
 - Interactive/General Purpose
 - Real-Time
- Scheduling algorithms attempt to maximize one or more “measurements”

Scheduling Algorithms

- For each algorithm discussed, we will assume we know the amount of time the CPU burst will last
 - Time is defined in clock cycles:
 - Faster computers have smaller clock cycles, but our numbers will remain the same
- Non-Preemptive: FCFS, SJF,
- Preemptive: Least Left First (LLF), Round Robin, Priority, Lottery

Summary

- ❑ Scheduling Algorithms are either Pre-emptive or non-preemptive
 - ❑ Each favor a particular set of criteria:
 - ❑ CPU Utilization, Wait Time, Responsiveness, etc
 - ❑ Implementation Concerns: **Speed**
 - ❑ Evaluate through simulation and observation
-


Multiprocessor Scheduling

- Goal: Load Sharing
- Assumptions: Homogenous set of CPU's
 - Note - CPU's may have dedicated bus to specific I/O (heterogeneous)

Asymmetric Multiprocessing ← Easy, but wasteful

Most efficient → SMP: Symmetric Multiprocessing

Processor Affinity

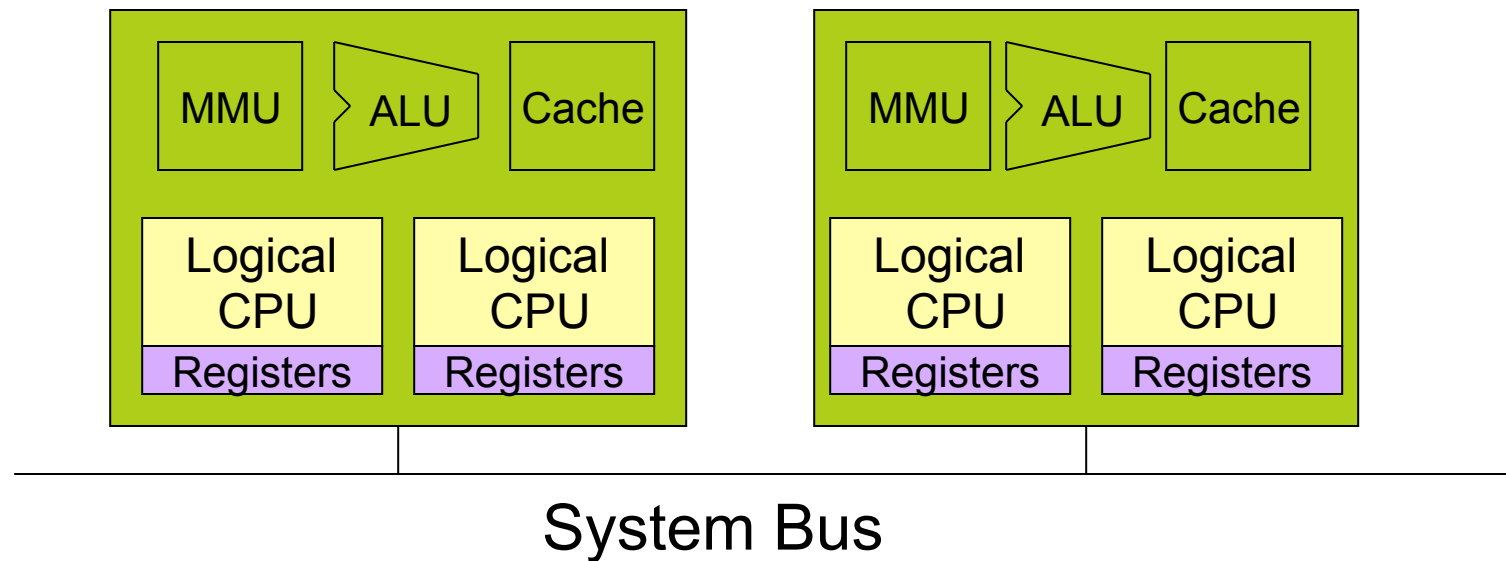


- *Process Migration*: Time Consuming
- *Processor Affinity*: Attempt to keep a process on the same CPU throughout its lifetime

Load Balancing

- ▣ When each CPU has separate ready queues, must prevent one from being overloaded
 - ▣ Push Migration
 - ▣ Pull Migration
 - ▣ How does this affect Processor Affinity?
 - ▣ Is this necessary when a common ready queue is used?
-

Symmetric Multi-threading



- ▣ Multiple CPU's are *emulated* on one chip.
- ▣ OS may/may not be aware.
 - ▣ Performance improved if made aware

Scheduling: Sun Solaris

- ▣ Kernel Threads
- ▣ Priority Based Scheduling:
 - ▣ Real-time, System, Time Sharing, Interactive
 - ▣ Within each priority, different scheduling algorithm
- ▣ Time Sharing and Interactive have dynamically changing priorities
 - ▣ Varies between 0 - 60 (low to high)

Changing Priorities

			Priority after...	
	Priority	Quantum Time	Quantum Expiration	Sleep/IO Completion
Lowest	0	200	0	50
	10	160	0	51
	20	120	15	52
	35	80	25	54
	50	40	40	58
Highest	59	20	49	59

What type of process frequently completes its quantum?

Scheduling: Windows

- Kernel Threads
 - Preemptive - Priority Scheduling
 - Highest priority thread **always** runs
 - Context Switch when...
 - Higher Priority thread is ready
 - Thread terminates
 - time quantum expires
 - blocking call issued

Priority Levels

- 32 Level priority
 - 2 Classes: variable and real-time
 - Variable: 1-15
 - Real-time: 16-31
 - Memory management: 0

Windows Priority Classes

Priorities

Relative Priorities

	real-time	high	above normal	normal	below normal	idle
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1

Priority lowered after preemption

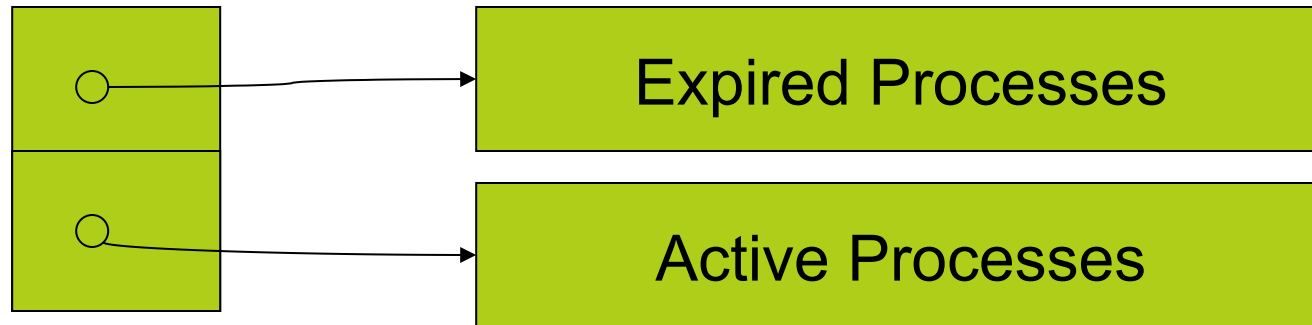
Priority elevated after I/O completion

Linux Scheduling

- Pre-emptive, Priority Based
 - Classes:
 - Real-time (0-99)
 - Nice (100-140)
 - Higher Priority gets **longer** quantum
 - Low priority - 10ms
 - High priority - 200ms
 - Notion of “expired” versus “active”

Split Ready Queue

“Ready Queue”



- ❑ Only Active Processes will run
- ❑ When a process uses its quantum, marked expired
- ❑ When active queue empties, pointers are swapped and process repeats