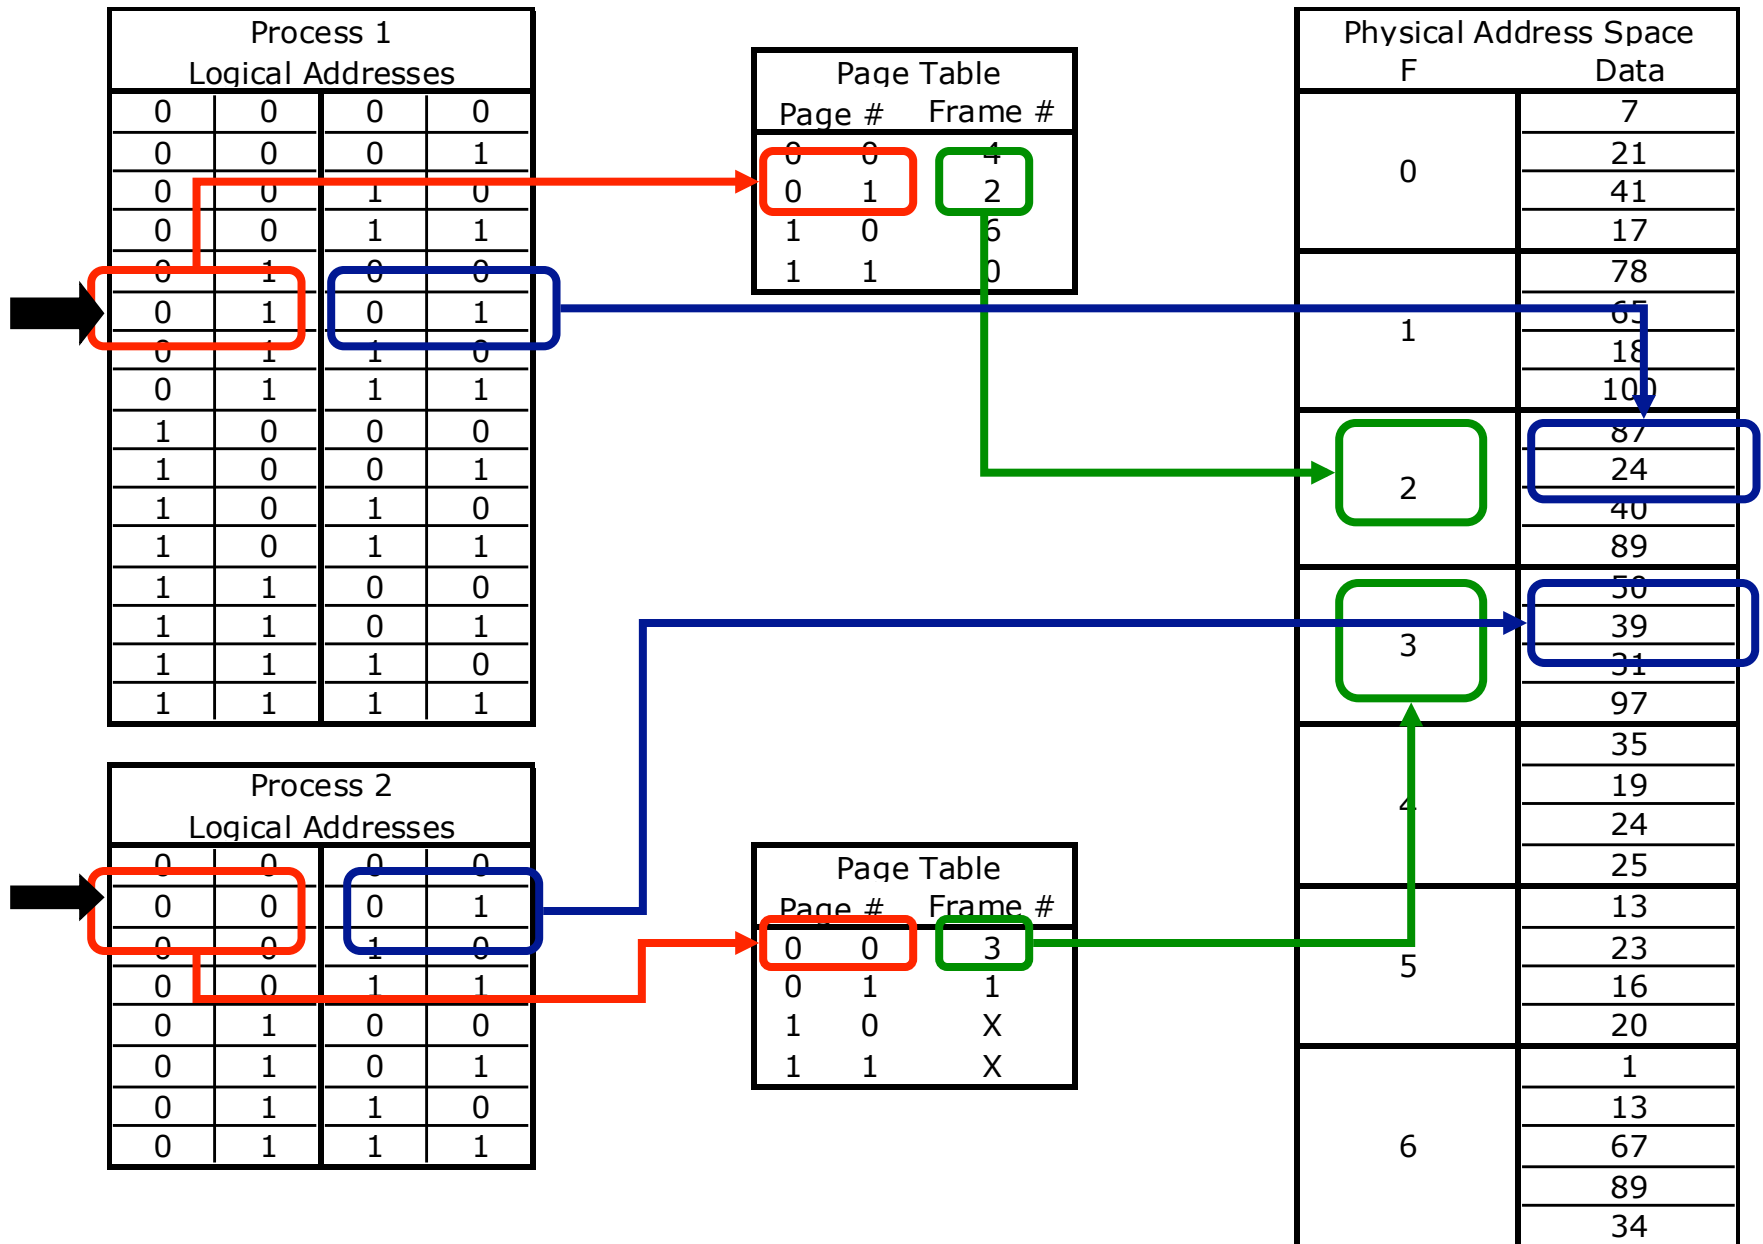


---

# Memory Management Virtual Memory

Chapter 8-9

---



# Paging in the real world

- ▣ 1 page table per process.
  - ▣ Page table may too big for **contiguous allocation**
  - ▣ What did we do the first time to solve contiguous allocation?

Page the Page Table too..

# More paging issues

- Even with multi-level paging strategies, page table(s) still quite large.
- One page table for each *logical* address space (1 per process).

Instead, one page table for *physical* addresses (frames).  
Inverted Page Table

# Virtual Memory

- Paging does not solve all our problems...

Remove the restriction of having an entire process in memory...

Program memory usage tends to be sparse...

This creates an “unlimited” address space.

# Page Fault

- ❑ Step 1: Translate Logical Address to determine page # & consult page table
- ❑ Step 2: If page is invalid, trap to OS
- ❑ Step 3: OS locates page on backing store
- ❑ Step 4: Find free frame in memory and copy page from backing store
- ❑ Step 5: Update page table
- ❑ Step 6: Restart program that initiated the request

# Demand paging

- We must decide *which* pages to bring into memory when program begins:
  - Pure Demand Paging: Only **page** pages when requested (when process begins, it has only one page)
  - Additional pages are only brought in when used
  - Lazy Swapping / Lazy Pager

---

# Page Replacement

- To implement demand paging, we need two things:
  - Frame Allocation Algorithm
  - Page Replacement Algorithm

**Overall Goal: Reduce Page Faults**

---



# Frame Allocation

- We'll get back to this... but
  - We'll be selecting a # of frames to allocate a specific process.
  - No matter what, more frame generally means less page faults (on average)
    - There are some anomalous circumstances where this isn't the case, for **every** situation

# Page Replacement

■ On a page fault, we must find a free frame in memory - what if its full?

- Select a ***victim*** frame to remove from memory to make room (put on backing store)
  - Page fault time doubles... (16,000,000ns)
  - Always copy back to backing store?

---

# Page Replacement Analysis

- To simulate a process running, we list the sequence of memory addresses:
    - We can keep track of which memory locations result in a page fault
    - This is called a *reference string*
  - We can simplify this to only list the page numbers to be accessed.
-

# Reference Strings

0100, 0432,  
0101, 0612, 0102, 0103, 0104,  
0101, 0611, 0102, 0103, 0104,  
0101, 0610, 0102, 0103, 0104,  
0101, 0609, 0102, 0105

If Page Size is 100 Bytes:

1, 4, 1, 6, 1, 6, 1, 6, 1, 6, 1

What is this  
code probably  
doing?

How many page  
faults if we have 1  
frame?

# Page Replacement Algorithms

■ 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

■ 3 Frames

- FIFO: Replace Oldest page
- Optimal: Replace page that won't be used for the longest time
- LRU: Replace the least recently used
- Second-Chance FIFO

# LRU Implementation

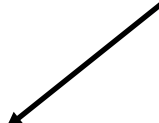
- Typically we don't keep track of exact timing...
  - Reference bits -> 00000
  - On each page access, set leftmost bit to 1
    - 10000
  - At given interval, shift **all page's** string right
    - 01000
- To find least recently used, find smallest number (likely to be many ties --- ok).
- To some extent, second chance FIFO is just LRU with one bit.

# Thrashing

- Very high rates of paging is **very** bad

- Typically called “thrashing”

What must this do to the  
page fault rate?



- Causes:

- CPU tries to allocate as many processes as possible to maintain CPU utilization

- Page Faults reduce CPU Utilization

- CPU responds by allocating even more processes...

# Working Set

- Programs exhibit locality -
  - They use the same group of memory addresses repeatedly for “relatively” long time periods
- Working Set: Pages in which a high percentage (90-99%) of memory addresses reside on
  - Working Set holds for some time period:  $\Delta$



# Avoiding Thrashing

- Approximate Working Set
- OS must allocate enough frames for a program to fit its current working set
  - If it cannot, must suspend the process
- Working Set leads to periodic page fault surges...

# Program Structure

```
int data[128][128];
```

```
for ( i = 0; i < 128; i++)  
    for ( j = 0; j < 128; j++)  
        data[i][j] = 0;
```

Page size = 128 integers (1024 bytes)

# Page Size?

- Large page size means...
  - Small page table
  - More internal fragmentation
  - Less IO time (more efficient)
  - Extends TLB reach
- Small page size means...
  - Larger page table
  - Less internal fragmentation
  - More IO time (less efficient)
  - Lower TLB reach

**In general, grow page sizes based on how much memory you have...**

# Exam 3 – Next Week

- Covers Memory Management basics
  - Logical vs. Physical Addresses
  - Compile / Load / Runtime Binding
  - Standard paging
    - Computing page size, # of pages based on address scheme
    - Resolving pages to frames
- Advanced Paging
  - Multi-Level Paging
  - TLB
  - Inverted Page Table

# Exam 3 – Next Week

## ■ Virtual Memory

- Understand each step towards resolving a page fault
- Be able to compute reference string
- Explain role of backing store
- Calculate # of page faults based on replacement strategy

## ■ Frame allocation

- Explain Working Set